





Ganeti Advanced Workshop

Ganeti Core Team - Google
GanetiCon 2013 - 3 Sept 2013



Monitoring Daemon

Michele Tartara <mtartara@google.com>

Once upon a time...

- there was no monitoring support in ganeti
- monitoring the status of the cluster was hard
 - Healthy?
 - Unhealthy?
- difficult to correlate hardware info with instances info (for an observer outside the cluster)
 - Instance <--> Logical volume ?
 - Instance <--> DRBD volume ?



What is the monitoring daemon?

Provides information:

- about the cluster state
- about the cluster health
 - automatically computed
- live
- read-only



What is NOT the monitoring daemon?

- A general-purpose monitoring system.
- Not meant to compete with Nagios, Pacemaker...
 - Integrate with them!
 - Provide them with easily parsable internal information

It just aims to monitor Ganeti and the related parts of the system



How is the monitoring daemon?

- HTTP daemon
- Replying to REST-like queries
 - Actually, GET only
- Providing JSON replies
 - Easy to parse in any language
 - Already used in all the rest of Ganeti
- Optional
- Dependent on Confd
- Built upon Haskell's Snap library



Where is the monitoring daemon?

- Running on every node
- Not:
 - Only master-candidates
 - Only VM-enabled



What info does it provide?

Now:

- instance status (Xen only)
- diskstats information
- LVM logical volumes information
- DRBD status information
- Node OS CPU load average

Soon(-ish):

- instance status (KVM)
- Ganeti daemons status
- Hypervisor resources
- Node OS resources report



Data collectors

- provide data to the deamon
- one collector, one report
- one collector, one topic



Two kinds of collectors (I)

- Performance reporting
 - only provide data "as is"
 - no interpretation



Two kinds of collectors (II)

- Status reporting
 - provide status evaluation
 - healthy
 - being auto-fixed
 - unknown
 - broken
 - hide data
 - informed status decisions require deep internals knowledge
 - to prevent meddling when auto-fixing
 - verbose mode
 - Not implemented yet
 - currently always-on



The query response

- JSON
- one report per collector, in a list

```
[  
  {  
    ... collector report  
  },  
  {  
    ... collector report  
  }  
]
```

JSON

- common structure for all the reports
- specific fields for each collector



The report format (I)

JSON

```
{
  "name" : "TheCollectorIdentifier",
  "version" : "1.2",
  "format_version" : 1,
  "timestamp" : 1351607182000000000,
  "category" : null,
  "kind" : 0,
  "data" : { "plugin_specific_data" : "go_here" }
}
```

- **name**: the name of the plugin. Unique string.
- **version**: the version of the plugin. A string.
- **format_version**: the version of the data format of the plugin. Incremental integer.
- **timestamp**: when the report was produced. Nanoseconds. Can be zero-padded.



The report format (II)

JSON

```
{
  "name" : "TheCollectorIdentifier",
  "version" : "1.2",
  "format_version" : 1,
  "timestamp" : 1351607182000000000,
  "category" : null,
  "kind" : 0,
  "data" : { "plugin_specific_data" : "go_here" }
}
```

- category: the category of the collector
 - storage, instance, daemon, hypervisor, "null"
 - can define a minimum set of prescribed fields
- kind: the kind of the collector
 - performance reporting (kind = 0)
 - status reporting (kind = 1, more to come)



The report format (III)

JSON

```
{  
  "name" : "TheCollectorIdentifier",  
  "version" : "1.2",  
  "format_version" : 1,  
  "timestamp" : 1351607182000000000,  
  "category" : null,  
  "kind" : 0,  
  "data" : { "plugin_specific_data" : "go_here" }  
}
```

- data: the collected data
 - free format
 - restrictions introduced by category and kind



Status reporting collectors: report

They introduce a mandatory part inside the data section.

```
"data" : {  
  ...  
  "status" : {  
    "code" : <value>  
    "message: "some summary goes here"  
  }  
}
```

JSON

- <value>: by increasing criticality level
 - 0: working as intended
 - 1: temporarily wrong. Being auto-repaired
 - 2: unknown. Potentially dangerous state
 - 4: problems. External intervention required



Status reporting collectors: report

They introduce a mandatory part inside the `data` section.

```
"data" : {  
  ...  
  "status" : {  
    "code" : <value>  
    "message: "some summary goes here"  
  }  
}
```

JSON

- `message`:
 - to better explain the reason of the status
 - optional (empty string) for codes 0 (ok) and 1 (auto-repair)
 - why could not be determined? (code 2, unknown)
 - what is wrong? (code 4, problems)



More goodies!

Enter mon-collector:

- quick 'n dirty CLI tool
- same collectors, same format
 - locally
- for quick checks by the sysadmin
- for local scripting
- shared code, different executables
 - works even if the daemon is not running
 - useful for offline-testing
- `$PREFIX/lib/ganeti/mon-collector <name>`
- `$PREFIX/lib/ganeti/tools/fmtjson`



What's where?

The development happens over time. Not everything is in every version.

- 2.7:
 - mon-collector
 - DRBD data collector
- 2.8:
 - monitoring daemon
- 2.9:
 - logical volumes collector
 - instance status (XEN) collector
 - /proc/diskstats
- 2.10:
 - CPU load (/proc/stat) collector. Thanks Spyros!



How to use the daemon?

- Accepts HTTP connections on `node.example.com:1815`
 - Not authenticated: read only
 - Just firewall, or bind on local address only
- GET requests to specific addresses
- Each address returns different info according to the API



The daemon API (I)

"Daemon, daemon on the port, would you send me a report?"

- `/`
return the list of supported protocol version numbers
- `/1`
the root of protocol version 1. Just returns `null`
- `/1/list/collectors`
list of (kind, category, name) tuples, representing all the collectors

The daemon API (II)

"Daemon, daemon on the port, would you send me a report?"

- `/1/report/all`
list of reports, one for each collector in the system
 - Will support `verbose=1`
- `/1/report/[category]/[collector_name]`
report produced by `[collector_name]` belonging to `[category]`



Collector categories

Storage

- Gather data about the storage subsystem
- Different levels of granularity and abstraction
 - Physical disks, partitions, LVs, ...
- Always possible to trace back to the instance
 - To find out performance problems
 - Instance directly provided whenever possible
 - "References" between levels elsewhere
 - device name, LV name, ...
- No common fields



Collector categories

Hypervisor

- Hypervisor's view of system resources
- No such collector, so quite undefined
- Status reporting / Performance reporting?
- Fields?
 - Free/used memory, #CPUs, CPU average load, ...

Collector categories

Daemon

- Gather data about Ganeti's own daemons
- Help identifying memory leaks, crashes, high resource utilization, ...
- Status reporting collectors (`kind = 1`)
- One collector per daemon
- Common fields in verbose mode:
 - `memory`
 - `uptime`
 - `cpu_usage` (percentage)



Collector categories

Instance

- Status reporting collectors (`kind = 1`)
- Reports a global status, and a per-instance status
- List of instances, with hypervisor-independent fields
 - `name`
 - `uuid`
 - `admin_state`
 - `actual_state`
 - `uptime`
 - `mtime`
 - `state_reason`



Stateful data collectors (2.10)

Again, thanks Spyros

- Stateless collectors
 - Traditional ones: data collected at invocation time
- Stateful collectors
 - Collection function
 - Collects the data
 - Run regularly by the monitoring daemon
 - Stores data in the daemon itself (memory, for now. Collectors, behave!)
 - Daemon-wide constant collection timer
 - Reporting function
 - Receives the collected data
 - Elaborates and prints them
 - Can collect more



Conclusion

- More details and complete list of fields of the collectors in the design doc: `doc/design-monitoring-agent.rst`
- Future work:
 - Plugin system
 - KVM instance status collector
 - More collectors
 - Per-collector collection function timer





The reason trail

Michele Tartara <mtartara@google.com>

The reason trail

- Initially required for the instance status (Xen) collector
 - Why did the instance last change its status?
 - Not just a message, but a complete track of what happened
 - Decisions:
 - What format for expressing this?
 - Where to store the information?



What format?

List of triples (source, reason, timestamp)

```
[("user", "Cleanup of unused instances", 1363088484000000000),  
 ("gnt:client:gnt-instance", "stop", 1363088484020000000),  
 ("gnt:opcode:shutdown", "job=1234;index=0", 1363088484026000000),  
 ("gnt:daemon:noded:shutdown", "", 1363088484135000000)]
```

PYTHON

- **source**: the entity deciding to perform/forward the command. Free form, but the **gnt:** prefix is reserved
- **reason**: why the entity decided to perform the operation
- **timestamp**: timestamp since epoch, in nanoseconds



Where is it?

- Inside every opcode
- `op["reason"]` field
- Visible with `gnt-job info`

```
root@node1:~# gnt-job info 4
```

```
Job ID: 4
```

```
...
```

```
Opcodes:
```

```
  OP_NODE_ADD
```

```
...
```

```
Input fields:
```

```
...
```

```
  reason: ['gnt:client:gnt-node', 'add', 1377589019102071040],  
          ['gnt:opcode:op_node_add', 'job=4;index=0', 1377589019106505984]
```

How is it generated?

- Automatically, from RAPI/CLI down to opcode level
- Before opcode generation:
 - User message (now):
 - CLI: `--reason`
 - RAPI: `reason` parameter added to the request
 - Previous trail (future, if useful)
- After opcode's job execution:
 - Specialized usages and manual implementations
 - Instance state change reason (start, stop, reboot. Serialized on file)



Conclusion

- Since Ganeti 2.8!
- More information available in the design doc: `doc/design-reason-trail.rst`
- Future work:
 - Accept an initial trail as input





Configuration Daemon (ConfD)

Michele Tartara <mtartara@google.com>

Once upon a t ...

For $t < 2.1$

- Configuration only available on master candidates
- Few selected values replicated with Ssconf
 - Small pieces of config in text files on all the nodes
 - Doesn't scale
- Need a way to access config from other nodes
 - Scalable
 - No single point of failure (so, no RAPI)



Enters ConfD

- Provides information from `config.data`
- Read-only
- Distributed
 - Multiple daemons running on master candidates
 - Accessible from all the nodes through ConfD protocol
 - Resilient to failures
- Optional



What info does it provide?

Replies to simple queries:

- Ping
- Master IP
- Node role
- Node primary IP
- Master candidates primary IPs
- Instance IPs
- Node primary IP from Instance primary IP
- Node DRBD minors
- Node instances



ConfD protocol


General description

- UDP (port 1814)
- keyed-Hash Message Authentication Code (HMAC) authentication
 - Pre-shared, cluster wide key
 - Generated at cluster-init
 - Root-only readable
- Timestamp
 - Checked (± 2.5 mins) to prevent replay attacks
 - Used as HMAC salt
- Queries made to any subset of master candidates
- Timeout
- Maximum number of expected replies




Confd protocol

Request/Reply



Confd protocol

Request/Reply



ConfD protocol

Request

```
plj0{  
  "msg": "{ \"type\": 1,  
            \"rsalt\": \"9aa6ce92-8336-11de-af38-001d093e835f\",  
            \"protocol\": 1,  
            \"query\": \"node1.example.com\" } \n",  
  "salt": "1249637704",  
  "hmac": "4a4139b2c3c5921f7e439469a0a45ad200aead0f"  
}
```

CONFID

- `plj0`: fourcc detailing the message content (PLain Json 0)
- `hmac`: HMAC signature of salt+msg with the cluster hmac key



ConfD protocol

Request

```
plj0{  
  "msg": "{\\"type\\": 1,  
    \\"rsalt\\": \\"9aa6ce92-8336-11de-af38-001d093e835f\\",  
    \\"protocol\\": 1,  
    \\"query\\": \\"node1.example.com\\"}\\"\\n",  
  "salt": "1249637704",  
  "hmac": "4a4139b2c3c5921f7e439469a0a45ad200aead0f"  
}
```

CONFID

- msg: JSON-encoded query
 - protocol: ConfD protocol version (=1)
 - type: What to ask for (CONFID_REQ_* constants)
 - query: additional parameters
- rsalt: response salt == UUID identifying the request



ConfD protocol

Reply

```
plj0{  
  "msg": "{\\"status\\": 0,  
          \\"answer\\": 0,  
          \\"serial\\": 42,  
          \\"protocol\\": 1}\\n",  
  "salt": "9aa6ce92-8336-11de-af38-001d093e835f",  
  "hmac": "aaecccc0dff9328fdf7967cb600b6a80a6a9332af"  
}
```

CONFD

- `salt`: the rsalt of the query
- `hmac`: hmac signature of salt+msg



ConfD protocol

Reply

```
plj0{  
  "msg": "{\\"status\\": 0,  
          \\"answer\\": 0,  
          \\"serial\\": 42,  
          \\"protocol\\": 1}\\n",  
  "salt": "9aa6ce92-8336-11de-af38-001d093e835f",  
  "hmac": "aaecccc0dff9328fdf7967cb600b6a80a6a9332af"  
}
```

CONFD

- msg: JSON-encoded answer
 - protocol: protocol version (=1)
 - status: 0=ok; 1=error
 - answer: query-specific reply
- serial: version of config.data



Ready-made clients

The protocol is simple, but clients are simpler

- Ready to use ConfD clients
 - Python
 - `lib/confd/client.py`
 - Haskell
 - Since Ganeti 2.7
 - `src/Ganeti/ConfD/Client.hs`
 - `src/Ganeti/ConfD/ClientFunctions.hs`



Expanding ConfD capabilities

- Currently not so many queries are supported
- Easy to add new ones
 - Just add a new query type in the constants list
 - ...and extend the `buildResponse` function
(`src/Ganeti/ConfD/Server.hs` to reply to it in the appropriate way



Conclusion

- More info in `doc/design-2.1.rst`
- Future work
 - More queries can be easily added as needed
 - Management of the configuration (on master) moved to a separate daemon from masterd





Autorepair (harep)

Michele Tartara <mtartara@google.com>

Before Ganeti 2.8

No self-repair

- DRBD instance is broken
 - manually fail it over
 - trigger a disk replacement
- Plain instance is broken
 - Manually recreate disk(s) and reinstall



Harep

- The Ganeti autorepair tool
- Available since Ganeti 2.8
- Meant to be run regularly using cron
- Admin can allow/disallow specific repairs



Controlling autorepair

- Harep is controlled through tags
- `ganeti:watcher:autorepair:<type>`
 - `instance/nodegroup/cluster`
 - What kind of repair allowed? (Sorted, more risky includes less risky)
 - `fix-storage`: disk replacement or fix the backend without affecting the instance itself (broken drbd secondary)
 - `migrate`: allow instance migration
 - `failover`: allow instance reboot on the secondary
 - `reinstall`: allow disks to be recreated and the instance to be reinstalled



Risks

- `fix-storage`: data loss if something is wrong on the primary but the secondary was somehow recoverable
- `migrate`: can cause instance crash (bugs)
- `failover`: loses the running state
- `reinstall`: data loss



Managing authorization conflicts

What if multiple autorepair tags act on an instance?

- In the same object: the least destructive takes precedence
- Across objects: the nearest tag wins
- Example:
 - cluster with I1 and I2
 - I1 has failover, the cluster has fix-storage and reinstall
 - Result: I1 failover, I2 fix-storage



Preventing autorepair

- Blocking a few repairs is easier than changing all the enabled ones
- `repair:suspended`
 - prevents an instance from being touched
 - can specify an expiration timestamp



How does it work?

- Multiple states for instances
 - Healthy
 - Suspended
 - Needs repair, repair disallowed
 - Pending repair
 - Failed
- Every run of harep
 - updates the tags
 - submits jobs



The result

`ganeti:watcher:autorepair:result:<type>:<id>:<timestamp>:
<result>:<jobs>`

- A `autorepair:result` tag is left on the repaired instance
- `<repair>`
 - success
 - failure
 - enoperm (=blocked by policies)

Conclusion

- More info `doc/design-autorepair.rst`
 - Includes detailed description of all the intermediate tags used internally





Cross-cluster instance migration

Guido Trotter <ultrotter@google.com>

Introduction

Instances can be moved between clusters that share a common secret.

- Operation available via the CLI or RAPI
- CLI tool uses RAPI, and can be seen as an example
- Data is transferred directly between clusters



Setup

Setup common secret and RAPI authentication

```
ssh root@cluster1 --> root@cluster1:~#  
gnt-cluster renew-crypto --new-cluster-domain-secret  
cat > /var/lib/ganeti/rapi/users <<EOF  
mover testpwd write  
EOF
```

```
# copy /var/lib/ganeti/cluster-domain-secret to the second cluster
```

```
ssh root@cluster2 --> root@cluster2:~#  
gnt-cluster renew-crypto --cluster-domain-secret=path_to_domain_secret  
# rapi access can be the same or different. in production use hashed passwords.  
cat > /var/lib/ganeti/rapi/users <<EOF  
mover testpwd write  
EOF
```

BASH



Execute move

Can be run on a third party machine

```
PWDFILE=$(mktemp)
echo testpwd > $PWDFILE

# Note: --dst-* defaults to --src-* if not specified
/usr/lib/ganeti/tools/move-instance --verbose \
  --src-ca-file=rapi.pem --src-username=mover \
  --src-password-file=$PWDFILE \
  [--dest-instance-name=new_name --net=0:mac=generate] \
  --iallocator=hail cluster1 cluster2 instance.example.com
```

BASH

Bugs:

- Either --iallocator or nodes must be specified manually
- Move is slower than it ought to be





hspace

Klaus Aehlig <aehlig@google.com>

Introduction

Capacity planning

- How many more instances can I add to my cluster?
- Which resource will I run out first?

So simulate sequentially adding new machines

- until we run out of resources
- allocation done as with hail
- start with maximal size of an instance (as allowed by the policy)
- reduce size if we hit the limit for one resource



On a live cluster

Use Luxi backend to get live cluster data

```
# hspace -L
```

BASH

The cluster has 3 nodes and the following resources:

MEM 196569, DSK 10215744, CPU 72, VCPU 288.

There are 2 initial instances on the cluster.

Tiered (initial size) instance spec is:

MEM 1024, DSK 1048576, CPU 8, using disk template 'drbd'.

Tiered allocation results:

- 4 instances of spec MEM 1024, DSK 1048576, CPU 8
- 2 instances of spec MEM 1024, DSK 258304, CPU 8
- most likely failure reason: FailDisk
- initial cluster score: 1.92199260
- final cluster score: 2.03107472
- memory usage efficiency: 3.26%
- disk usage efficiency: 92.27%
- vcpu usage efficiency: 18.40%

[...]

The simulation backend

One of the lesser known backends (hspace and hail)
Mainly for cluster planning

- Simulates an empty cluster with given data
- Format
 - allocation policy (p=preferred, a=last resort, u=unallocatable)
 - number of nodes (in this group)
 - disk space per node (in MiB)
 - ram (in MiB)
 - number of physikal CPUs
- use --simulate several times for more node groups



Planning a cluster

What if I bought 10 times more disks?

BASH

```
$ hspace --simulate=p,3,34052480,65523,24 \  
> --disk-template=drbd --tiered-alloc=1048576,1024,8  
The cluster has 3 nodes and the following resources:  
  MEM 196569, DSK 102157440, CPU 72, VCPU 288.  
There are no initial instances on the cluster.  
Tiered (initial size) instance spec is:  
  MEM 1024, DSK 1048576, CPU 8, using disk template 'drbd'.  
Tiered allocation results:  
  - 33 instances of spec MEM 1024, DSK 1048576, CPU 8  
  - 3 instances of spec MEM 1024, DSK 1048576, CPU 7  
  - most likely failure reason: FailCPU  
  - initial cluster score: 0.000000000  
  - final cluster score: 0.000000000  
  - memory usage efficiency: 18.75%  
  - disk usage efficiency: 73.90%  
  - vcpu usage efficiency: 100.00%  
[...]
```



hroller

Klaus Aehlig <aehlig@google.com>

Introduction

When rebooting all nodes (e.g., kernel update), there are several things to take care of.

- Don't reboot primary and secondary at the same time. Machine/disks might not come back after reboot.
- When doing live migration, have enough memory. No two nodes with primaries, that have the same secondary.
- When fully evacuating, plan for disk space.



Default

hroller suggests groups of nodes to be rebooted together.
By default, plan for live migration.

```
# hroller -L  
'Node Reboot Groups'  
node-00,node-10,node-20,node-30  
node-01,node-11,node-21,node-31
```

BASH

Also possible to only avoid primary/secondary reboots (--offline-maintenance) or to plan complete node evacuation (--full-evacuation).

```
# hroller -L --full-evacuation  
'Node Reboot Groups'  
node-01,node-11  
node-00,node-10  
node-20,node-30  
node-21,node-31
```

BASH

Moves

For the full evacuation, moves can also be shown (--print-moves). Typically, together with --one-step-only.

```
# hroller -L --full-evacuation --print-moves --one-step-only
```

BASH

```
'First Reboot Group'
```

```
node-01
```

```
node-11
```

```
inst-00 node-00 node-20
```

```
inst-00 node-00 node-10
```

```
inst-10 node-10 node-21
```

```
inst-11 node-10 node-00
```



Tags

Nodes to be considered can also be selected by tags.
This allows reboots interleaved with other operations.

```
GROUP=`hroller --node-tags needsreboot --one-step-only --no-headers -L`  
for node in $GROUP; do gnt-node modify -D yes $node; done  
for node in $GROUP; do gnt-node migrate -f --submit $node; done  
# ... wait for migrate jobs to finish  
# reboot nodes in $GROUP  
# verify...  
for node in $GROUP; do gnt-node remove-tags $node needs-reboot; done  
for node in $GROUP; do gnt-node modify -D no $node; done  
hspace -L -X
```

BASH



Thank You!

Questions?

