

# Queue and RAPI Daemon Speed

Christos Stavrakakis

Greek Research & Technology Network

*cstavr@grnet.gr*

September 4, 2013





Question:



Question:

Say we want to create 100 instances...



Question:

Say we want to create 100 instances...  
at once...



## Question:

Say we want to create 100 instances...  
at once...  
How long do yo think this would take?

# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent

# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent
- Critical to keep request-response cycle low: enqueue jobs to Ganeti, and get the job ID



# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent
- Critical to keep request-response cycle low: enqueue jobs to Ganeti, and get the job ID
- **No polling!** Instead, receive asynchronous notifications from the Ganeti to track job progress

# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent
- Critical to keep request-response cycle low: enqueue jobs to Ganeti, and get the job ID
- **No polling!** Instead, receive asynchronous notifications from the Ganeti to track job progress
- Use Ganeti hooks?

# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent
- Critical to keep request-response cycle low: enqueue jobs to Ganeti, and get the job ID
- **No polling!** Instead, receive asynchronous notifications from the Ganeti to track job progress
- Use Ganeti hooks?
  - Post hooks run only on success. How to track errors?

# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent
- Critical to keep request-response cycle low: enqueue jobs to Ganeti, and get the job ID
- **No polling!** Instead, receive asynchronous notifications from the Ganeti to track job progress
- Use Ganeti hooks?
  - Post hooks run only on success. How to track errors?
  - Time overhead to total job execution time.

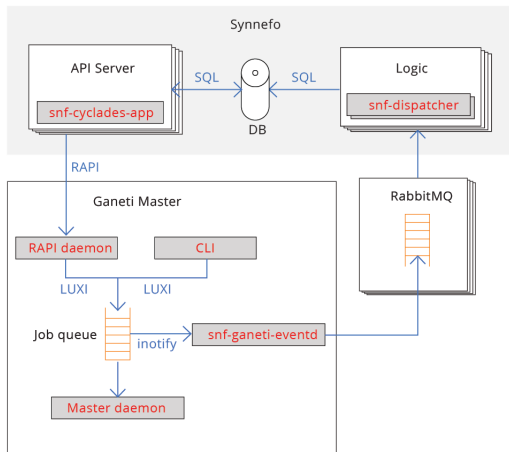
# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent
- Critical to keep request-response cycle low: enqueue jobs to Ganeti, and get the job ID
- **No polling!** Instead, receive asynchronous notifications from the Ganeti to track job progress
- Use Ganeti hooks?
  - Post hooks run only on success. How to track errors?
  - Time overhead to total job execution time.
- **inotify**: Notice changes to Ganeti job queue directory!

# How Synnefo uses Ganeti (1)

- Use Ganeti RAPI: no need for custom API and agent
- Critical to keep request-response cycle low: enqueue jobs to Ganeti, and get the job ID
- **No polling!** Instead, receive asynchronous notifications from the Ganeti to track job progress
- Use Ganeti hooks?
  - Post hooks run only on success. How to track errors?
  - Time overhead to total job execution time.
- **inotify**: Notice changes to Ganeti job queue directory!
- reconciliation: Query via RAPI to reconcile state of DB with Ganeti

## How Synnefo uses Ganeti (2)



Time to create those instances



Time to create those instances  
First count how long it takes to enqueue the jobs:

Time to create those instances

First count how long it takes to enqueue the jobs:

- Enqueue 1 job  $\rightarrow$  0.2s

## Time to create those instances

First count how long it takes to enqueue the jobs:

- Enqueue 1 job  $\rightarrow 0.2s$
- Enqueue jobs sequentially  $\rightarrow 10jobs/sec$

## Time to create those instances

First count how long it takes to enqueue the jobs:

- Enqueue 1 job  $\rightarrow 0.2s$
- Enqueue jobs sequentially  $\rightarrow 10jobs/sec$
- Enqueue jobs in parallel

## Time to create those instances

First count how long it takes to enqueue the jobs:

- Enqueue 1 job  $\rightarrow 0.2s$
- Enqueue jobs sequentially  $\rightarrow 10jobs/sec$
- Enqueue jobs in parallel

## Time to create those instances

First count how long it takes to enqueue the jobs:

- Enqueue 1 job  $\rightarrow 0.2s$
- Enqueue jobs sequentially  $\rightarrow 10jobs/sec$
- Enqueue jobs in parallel  $\rightarrow 10jobs/sec$

# Enqueue jobs

## GanetiRAPIClient

- Powered by PycURL → libcurl

# Enqueue jobs

## GanetiRAPIClient

- Powered by PycURL → libcurl
- Does not work well with gevent



# Enqueue jobs

## GanetiRAPIClient

- Powered by PycURL → libcurl
- Does not work well with gevent
- Trivial client implementation using Requests
  - Powered by urllib3
  - Greenlet-safe
  - HTTP connection pooling

# Enqueue jobs

## GanetiRAPIClient

- Powered by PycURL → `libcurl`
- Does not work well with `gevent`
- Trivial client implementation using `Requests`
  - Powered by `urllib3`
  - Greenlet-safe
  - HTTP connection pooling
- Enqueue 1 job → 0.2s
- Enqueue jobs sequentially → 10*jobs/sec*
- Enqueue jobs in parallel → 10*jobs/sec*
- Enqueue jobs with green client → 33*jobs/sec*

# Enqueue jobs

## GanetiRAPIClient

- Powered by PycURL → `libcurl`
- Does not work well with `gevent`
- Trivial client implementation using `Requests`
  - Powered by `urllib3`
  - Greenlet-safe
  - HTTP connection pooling
- Enqueue 1 job → 0.2s
- Enqueue jobs sequentially → 10*jobs/sec*
- Enqueue jobs in parallel → 10*jobs/sec*
- Enqueue jobs with green client → 33*jobs/sec*

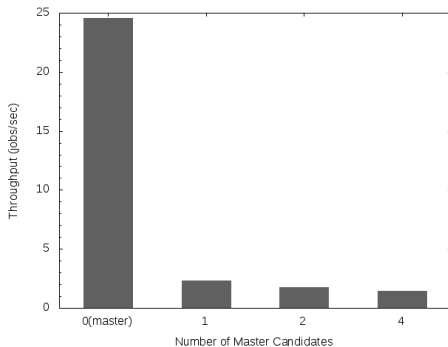
But, on a different Ganeti cluster enqueueing jobs was much slower. Why?

# Master Candidates

- Number of master candidates affects job submission rate
- Ganeti writes job to disk and then concurrently replicates it to master candidates


# Master Candidates

- Number of master candidates affects job submission rate
- Ganeti writes job to disk and then concurrently replicates it to master candidates



# Master Candidates

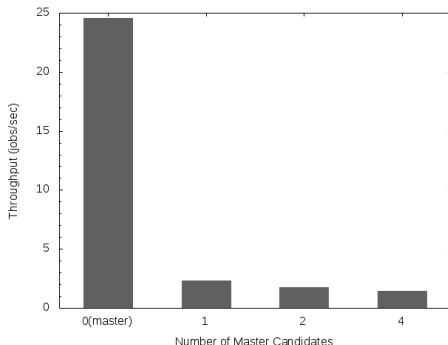
- Number of master candidates affects job submission rate
- Ganeti writes job to disk and then concurrently replicates it to master candidates



- With 1MCs  $\rightarrow \approx 7$  times slower

# Master Candidates


- Number of master candidates affects job submission rate
- Ganeti writes job to disk and then concurrently replicates it to master candidates



- With 1MCs  $\rightarrow \approx 7$  times slower
- $> 2$ MCs  $\rightarrow$  no significant performance drop

# Master Candidates

- Number of master candidates affects job submission rate
- Ganeti writes job to disk and then concurrently replicates it to master candidates

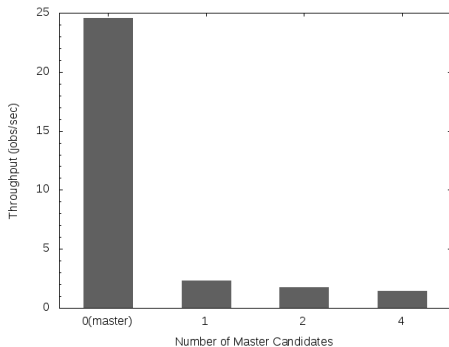


- With 1MCs  $\rightarrow \approx 7$  times slower
- $> 2$ MCs  $\rightarrow$  no significant performance drop
- Job queue directory in SSDs + MCs in non vm\_capable nodes



# Master Candidates

- Number of master candidates affects job submission rate
- Ganeti writes job to disk and then concurrently replicates it to master candidates



- With 1MCs  $\rightarrow \approx 7$  times slower
- $> 2MCs \rightarrow$  no significant performance drop
- Job queue directory in SSDs + MCs in non vm\_capable nodes
- Asynchronous replication?

# Job execution phases

Jobs enqueued! Awaiting for completion...

# Job execution phases

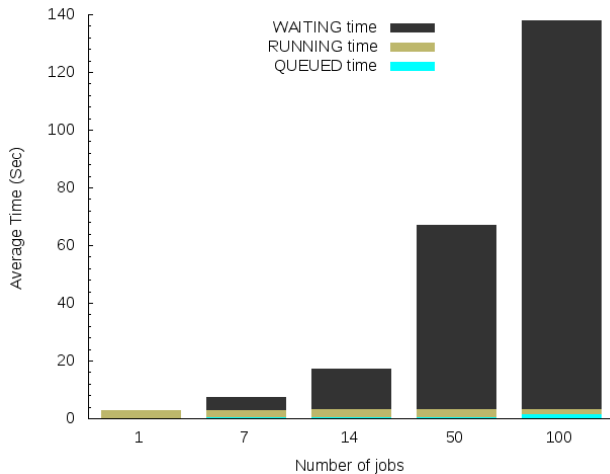
Jobs enqueued! Awaiting for completion...

- Used a Ganeti cluster of 7 vm\_capable nodes
- Minimized running time (no-install, 1G file disk)  $\approx 3\text{sec}$

# Job execution phases

Jobs enqueued! Awaiting for completion...

- Used a Ganeti cluster of 7 vm\_capable nodes
- Minimized running time (no-install, 1G file disk)  $\approx 3\text{sec}$



# Locking

- start/stop/modify/.. → lock instance

# Locking

- start/stop/modify/.. → lock instance
- remove/migrate/.. → lock primary/secondary node

# Locking

- start/stop/modify/.. → lock instance
- remove/migrate/.. → lock primary/secondary node
- create:
  - allocate instance → lock all nodes
  - afterwards → lock primary/secondary node
  - Use **opportunistic locking**( $\geq$  *Ganeti* 2.7)

# Locking

- start/stop/modify/.. → lock instance
- remove/migrate/.. → lock primary/secondary node
- create:
  - allocate instance → lock all nodes
  - afterwards → lock primary/secondary node
  - Use **opportunistic locking** ( $\geq$  *Ganeti* 2.7)
- verify cluster → lock all nodes
  - Use **nodegroups**



# Real-World Scenario

100 instances: 2 CPU, 4 GB RAM, 40GB Disk  
2 x Ganeti cluster of 7 nodes

# Real-World Scenario

100 instances: 2 CPU, 4 GB RAM, 40GB Disk  
2 x Ganeti cluster of 7 nodes

- drbd:  $\approx 14.3min$
- archipelago:  $\approx 4.1min$

- Based on "[RFC] Distributed Database Config Storage"

# Ongoing Work

- Based on "[RFC] Distributed Database Config Storage"
- Store job queue and config data in a distributed NoSQL database (CouchDB)

# Ongoing Work

- Based on "[RFC] Distributed Database Config Storage"
- Store job queue and config data in a distributed NoSQL database (CouchDB)
- config.data: JSON-serialized file
  - Is read once from file and then kept in memory
  - Is written to file in every update
  - Updating config file can be slow when the size of clusters increases
  - Why not break it into instances, nodes, ...

# Ongoing Work

- Based on "[RFC] Distributed Database Config Storage"
- Store job queue and config data in a distributed NoSQL database (CouchDB)
- config.data: JSON-serialized file
  - Is read once from file and then kept in memory
  - Is written to file in every update
  - Updating config file can be slow when the size of clusters increases
  - Why not break it into instances, nodes, ...
- Implemented as optional mode of operation

# Ongoing Work

- Based on "[RFC] Distributed Database Config Storage"
- Store job queue and config data in a distributed NoSQL database (CouchDB)
- config.data: JSON-serialized file
  - Is read once from file and then kept in memory
  - Is written to file in every update
  - Updating config file can be slow when the size of clusters increases
  - Why not break it into instances, nodes, ...
- Implemented as optional mode of operation
- Replication in DB, easier master failover, ...

Thank you.

Questions?