

ganeti.deb

Packaging Ganeti in Debian

Apollon Oikonomopoulos

`apoikos@debian.org`



GanetiCon 2015
15-17 Sep 2015 — Prague, CZ

Ganeti in Debian

- ▶ Team-maintained: `pkg-ganeti-devel@lists.alioth.debian.org`
- ▶ 3 team members (help appreciated :)
- ▶ Package source tracked in git
`http://anonscm.debian.org/gitweb/?p=pkg-ganeti/ganeti.git`



Release policy

- ▶ Latest stable release in unstable
 - ▶ Backport to jessie-backports as soon as it hits testing
- ▶ RCs usually in experimental
- ▶ PPA for Ubuntu LTS releases (currently 12.04 and 14.04)



News since GanetiCon1



Small bits

- ▶ Multiple versions support: tested heavily, no complaints so far (but must go through NEW for every minor release)
- ▶ 2.12 in Jessie: 2.12.0 updated to 2.12.4 via a stable update, most issues sorted out
- ▶ GHC 7.8 in unstable/testing: compatibility is fine
- ▶ GHC 7.10.2 in experimental: integer-gmp2 as default (?): check compatibility and re-evaluate SSL library choice.
- ▶ Switched the packaging repository to upstream-git: easier to cherry-pick patches (esp. autotools-related)



DEP-8 automatic as-installed tests

- ▶ Implemented automatic package tests using `vcluster` to run the QA suite.
- ▶ Test the actual built Debian package in a clean environment:
 - ▶ Useful to catch both packaging and upstream errors.
 - ▶ Runs with the "real" settings (most notably multiple user support).
- ▶ Requires `isolation-container` – currently does not run on `ci.debian.net`
- ▶ ... but it runs on my laptop before every upload :-)



Debian bug #783683

- ▶ cron's systemd unit has `KillMode=control-group`
- ▶ The watcher by default starts noded if it's not running → `ganeti-noded` ends up in `cron.service` cgroup.
- ▶ KVM processes spawned subsequently remain in `cron.service` cgroup.
- ▶ Restarting cron kills noded and all KVM processes in the cgroup.



Debian bug #783683

- ▶ cron's systemd unit has `KillMode=control-group`
- ▶ The watcher by default starts noded if it's not running → `ganeti-noded` ends up in `cron.service` cgroup.
- ▶ KVM processes spawned subsequently remain in `cron.service` cgroup.
- ▶ Restarting cron kills noded and all KVM processes in the cgroup.

KABOOM!



Debian bug #783683

- ▶ cron's systemd unit has `KillMode=control-group`
- ▶ The watcher by default starts noded if it's not running → `ganeti-noded` ends up in `cron.service` cgroup.
- ▶ KVM processes spawned subsequently remain in `cron.service` cgroup.
- ▶ Restarting cron kills noded and all KVM processes in the cgroup.

KABOOM!

- ▶ Although now fixed, there are some questions:
 - ▶ Do we need ganeti-watcher starting noded? How do you actually stop noded?
 - ▶ Should we coordinate with systemd and place instances under `machine.slice` (and depend on `python-dbus`)?



Issues with QEMU 2.1

- ▶ Two known bugs



Issues with QEMU 2.1

- ▶ Two known bugs
- ▶ #785557: Steal time overflow on live migration
 - ▶ Caused `perl` build failures on Debian's build's
 - ▶ May cripple the guest kernel's scheduler (3.16+ affected, possibly newer as well). The system reports 100% steal time.
 - ▶ Can be worked around: `cpu_type=qemu64,-kvm_steal_time`
 - ▶ Forwarded upstream, no reaction sofar



Issues with QEMU 2.1

- ▶ Two known bugs
- ▶ #785557: Steal time overflow on live migration
 - ▶ Caused `perl` build failures on Debian's build's
 - ▶ May cripple the guest kernel's scheduler (3.16+ affected, possibly newer as well). The system reports 100% steal time.
 - ▶ Can be worked around: `cpu_type=qemu64,-kvm_steal_time`
 - ▶ Forwarded upstream, no reaction so far
- ▶ #786789: VM hangs post-migration due to TSC jumps
 - ▶ Fixed in QEMU 2.1.3
 - ▶ The fix is not yet in Jessie → use 2.3 from sid/stretch or build your own packages.



- ▶ External repository for debian packages.
- ▶ Needed for 2.12+ on wheezy.
- ▶ Will probably offer long-ish term support for stable branches for multiple distributions (Wheezy and Jessie for now).
- ▶ amd64-only for now, needs proper build infrastructure.



Future work



ganeti-quickstart

- ▶ Goal: make it easy for our users to setup a small cluster



ganeti-quickstart

- ▶ Goal: make it easy for our users to setup a small cluster
- ▶ Use-cases: personal VM manager, sneak preview



ganeti-quickstart

- ▶ Goal: make it easy for our users to setup a small cluster
- ▶ Use-cases: personal VM manager, sneak preview
 - ▶ Single-node cluster
 - ▶ KVM (does not need reboot, setup is fairly easy)
 - ▶ Primarily file storage (LVM not guaranteed to be configured or have free extents)
 - ▶ Pre-configure ganeti-instance-debootstrap
- ▶ `ganeti-quickstart` package



Thank you!

Q&A

