

random thoughts on Ganeti operation

plain KVM/DRBD/bridged

Sascha Lucas

GISA GmbH Halle (Saale)

18th of June 2019 Ganeticon / Umeå

node OS: Ubuntu 16.04

- plain package 2.15.2-3 with some tweaks

- ▶ livemigration progress broken

```
--- /old/usr/share/ganeti/2.15/ganeti/hypervisor/hv_kvm/__init__.py
+++ /new/usr/share/ganeti/2.15/ganeti/hypervisor/hv_kvm/__init__.py
@@ -379,7 +379,7 @@
     _MIGRATION_STATUS_RE = re.compile(r"Migration\s+status:\s+(\w+)",
                                         re.M | re.I)

     _MIGRATION_PROGRESS_RE = \
-        re.compile(r"\s*transferred\s+ram:\s+(\?P<transferred>\d+)\s+kbytes\s*\n"
+        re.compile(r"\s*transferred\s+ram:\s+(\?P<transferred>\d+)\s+kbytes\s*\n.*"
+            r"\s*remaining\s+ram:\s+(\?P<remaining>\d+)\s+kbytes\s*\n"
+            r"\s*total\s+ram:\s+(\?P<total>\d+)\s+kbytes\s*\n", re.I)
```

- ▶ bash completion doesn't work out of the box

```
ln -s /usr/share/bash-completion/completions/ganeti /etc/bash_completion.d/
```

- ▶ ganeti-instance-debootstrap broken → use 0.16-3ubuntu1 from 18.04/bionic

- good coverage of canonical 5y supported packages: `command` `ubuntu-support-status`

Ganeti „distributed switch“

- setup node with `/etc/network/interfaces`

```
source /etc/network/interfaces.group1/*
```

- put VLAN and bridge setup into one file: `/e/n/i.group1/net100`

```
auto bond0.100
iface bond0.100 inet manual
    vlan-raw-device bond0
    mtu 9000

auto br100
iface br100 inet manual
    bridge_ports bond0.100
    bridge_fd 0
    bridge_stp off
    bridge_maxwait 1
    mtu 9000
```

- distribute config to the cluster:

```
gnt-cluster copyfile /etc/network/interfaces.group1/net100
```

- bring up the interface: `gnt-cluster command ifup bond0.100`¹
- but ... VLAN aware bridge would be simpler/better

¹ignore errors from nodes not belonging to node group 1

memory management

part 1

what is default in Ubuntu (I assume in Debian too)

- qemu enables KSM by default: `/etc/default/qemu-kvm`
- KSM merges across NUMA nodes: `/sys/kernel/mm/ksm/merge_across_nodes`
- THP is enabled per default for qemu (madvise):
`/sys/kernel/mm/transparent_hugepage/enabled`
- kernel NUMA balancer is enabled per default on multisocket systems

```
$ sysctl kernel.numa_balancing
kernel.numa_balancing = 1
```

All this scans your memory and possibly work against each other. This is probably not what NUMAists want.

memory management

part 2

numad to the rescue

- monitors your system and dynamically adjusts NUMA locality (memory and CPU-masks)
- sets THP_scan_sleep_ms from 10 to 1s
- manually disable qemu KSM
- manually disable kernel NUMA balancer

numad drawbacks

- sometimes too dynamic
- depends on the amount of CPU oversubscribe and instance CPU utilization
- NUMA-Node ping-pong → can trigger instance kernel watchdog on blocking CPU

memory management

numad and THP example

```
$ numastat -c qemu
Per-node process memory usage (in MBs)
PID           Node 0 Node 1 Total
-----
1003 (qemu-syste      8  32812 32819
1817 (qemu-syste      8  19339 19347
2556 (qemu-syste  4146      0  4146
3326 (qemu-syste  2094      0  2094
4502 (qemu-syste  2088      0  2088
5290 (qemu-syste      7  2097  2104
6050 (qemu-syste  2103      0  2103
6772 (qemu-syste  4121      0  4121
7492 (qemu-syste  6139      0  6139
8242 (qemu-syste  4108      0  4108
9003 (qemu-syste  4130      0  4130
9746 (qemu-syste  6166      0  6166
10511 (qemu-syst  4146      0  4146
12092 (qemu-syst  3420      0  3421
23702 (qemu-syst  7191      0  7191
28563 (qemu-syst      8  2449  2457
29938 (qemu-syst 16435      0 16436
30626 (qemu-syst      8  24537 24545
32157 (qemu-syst  8912      0  8912
-----
Total           75236  81235 156472
```

```
$ lscpu | grep NUMA
NUMA node(s):          2
NUMA node0 CPU(s):     0-7,16-23
NUMA node1 CPU(s):     8-15,24-31
```

```
$ taskset -cp 1003
pid 1003's current affinity list: 8-15,24-31
```

```
$ grep -i huge /proc/meminfo
AnonHugePages: 146348032 kB
```

post-copy migration

- memory write intensive VMs are hard to migrate

- ▶ even with high migration_bandwidth=1000
- ▶ high migration_downtime (>30 ms) may not be tolerable by instances
- ▶ ... endless copying memory

- solution: post-copy migration

- ▶ memory is copied after switching execution state from source to target node
- ▶ steps to use:

```
$ gnt-cluster modify -H kvm:migration_caps=postcopy-ram # (or x-postcopy-ram on qemu-2.5)
```

on instance source node (ideally after ≥ 1 cycle/100% of memory transfer)

```
$ echo "migrate_start_postcopy" | socat \  
STDIO UNIX-CONNECT:/var/run/ganeti/kvm-hypervisor/ctrl/some.vm.monitor
```

- ▶ migrate_start_postcopy command must timed right to not confuse Ganetis migration status pull (info migrate)
- ▶ in some development branch this feature was added to Ganeti. But it seems never released???

DRBD: simple, stupid

- can be called: software defined / distributed storage (hyper converged)
- needs little to no knowledge to be used by Ganeti
- however Ganeti has non optimal defaults
 - ▶ users struggle with static resync vs. dynamic resync controller (which one is active/to tune)
 - ▶ higher resync speeds (> 100 MB/s) needs larger buffers
i.e. $150\text{MB/s} \rightarrow \text{net-custom}='--\text{max-buffers } 8000 --\text{max-epoch-size } 8000'$
 - ▶ Ganeti's setting `c-plan-ahead=0` leads to skip DRBDs `c-plan-ahead`, which is 20 per default
DRBDs `c-plan-ahead != 0` enables the dynamic resync controller
forcefully disable with `disk-custom='--c-plan-ahead 0'`
 - ▶ are even Debian DSA's struggling? <https://dsa.debian.org/howto/install-ganeti/> (section DRBD optimization)
- split brain during live migration
 - ▶ DRBD is in dual primary during migration
 - ▶ split brain (standalone/standalone) will happen if DRBD gets disconnected during migration
 - ▶ happens when migration bandwidth saturates the link
 - ▶ qemu will finish migration, because it does not know anything about DRBD

hooks and tweaks

- growing a disk online: see <https://github.com/saschalucas/ganeti-hook-grow-disk>
- network anti spoofing
 - ▶ prevent MAC, ARP and IP spoofing
 - ▶ combination of up script (`$CONF_DIR/kvm-vif-bridge`) and in absence of down script `instance-stop.pre-d` hook
- I/O-limit your instances with `cgroup-v1/blkio-controller`
 - ▶ needs `DIRECTIO` (`cache=none`)
 - ▶ a single PV (major/minor) of the ganeti VG
 - ▶ example: limit 1000 write IOPS and 300MB/s read

```
mkdir /sys/fs/cgroup/blkio/some.vm
# here 8:16=/dev/sdb substitute with your major/minor number
echo "8:16 1000" > /sys/fs/cgroup/blkio/some.vm/blkio.throttle.write_iops_device
echo "8:16 314572800" > /sys/fs/cgroup/blkio/some.vm/blkio.throttle.read_bps_device
echo PID_OF_QEMU_INST > /sys/fs/cgroup/blkio/some.vm/cgroup.procs
```

- CPU fairness: i.e. 8 Core node, two instances: 4 vCPUs and 8 vCPUs
 - ▶ without cgroups: $4c=1/3$ and $8c=2/3$
 - ▶ with cgroups: $4c=1/2$ and $8c=1/2$

security

not enabled per default

- Want your customers share their data? No? Wiping disk is obligatory:
`--prealloc-wipe-disks=yes`
- KVM can use chroot: `--hypervisor-parameters=kvm:use_chroot=true`
- UID separation between VMs on the same node:
 - ▶ use <https://github.com/grnet/nss-uidpool> to get 100 UIDs without creating 100 users on each node
 - ▶ and `--uid-pool=10002-10100 --hypervisor-parameters=kvm:security_model=pool`

security

Spectre/MDS mitigation not enabled in qemu by default

- not enabled by your hardware: disable SMT
check `/sys/devices/system/cpu/vulnerabilities/mds`
- showing available CPU models: `qemu-system-x86_64 -cpu ?`
- chose a IBRS variant and enable this flags:
 - ▶ `+pcid`: mitigate the cost of the Meltdown
 - ▶ `+ssbd`: required to enable the CVE-2018-3639 fix
 - ▶ `+md-clear`: signal that host can mitigate MDS
 - ▶ `enforce`: don't start if the host can't fulfill the desired CPU type/flags
 - ▶ see also: <https://www.berrange.com/tags/ssbd/>
- testing the cluster for a common CPU the model:

```
$ gnt-cluster command 'qemu-system-x86_64 -cpu XXXXXXX-IBRS,+pcid,+ssbd,+md-clear,enforce \
-machine accel=kvm -nographic -nodefaults -boot c,reboot-timeout=1 -no-reboot'
```

- setting a cpu type

```
$ gnt-cluster modify -H kvm:cpu_type='XXXXXXX-IBRS\,+pcid\,+ssbd\,+md-clear\,enforce'
```

instance creation with ganeti-instance-debootstrap

preparation

- features
 - ▶ uses kernel and initrd from node (best when node and instance OS are same)
 - ▶ no kernel package inside instance (modules from initrd must be sufficient)
 - ▶ no boot loader (grub)
 - ▶ caches debootstrap in a tar file

- `/etc/default/ganeti-instance-debootstrap`

```
MIRROR="https://put.your-mirror.here/ubuntu"  
ARCH="amd64"  
SUITE="xenial"  
EXTRA_PKGS="acpid,ssh"  
COMPONENTS="main,universe"  
GENERATE_CACHE="yes"  
CLEAN_CACHE="14"
```

- `/etc/ganeti/instance-debootstrap/hooks/000interfaces`

- ▶ enhanced with instance network configuration WRT gnt-network
- ▶ see <https://github.com/ganeti/instance-debootstrap/pull/1>

instance creation with ganeti-instance-debootstrap

in action

```
gnt-instance add -t drbd --disk 0:size=1G --net 0:network=gruen97,ip=172.28.97.234\  
-H kvm:kernel_path=/vmlinuz,initrd_path=/initrd.img,root_path=/dev/vda1,\  
kernel_args='ro elevator=noop net.ifnames=0'\  
-B vcpus=1,memory=1G -o debootstrap+default --no-name-check --no-ip-check test.vm
```

make sure:

- specify an IP if you don't have DHCP
- ganeti-instance-debootstrap assumes eth0 as NIC name, so don't forget net.ifnames=0

very fast:

- 14s from command submission to instance startup
- 30s from command submission to instance first ping

ganeti-instance-debootstrap can be enhanced i.e. by

- using libguestfs (i.e. install kernel/grub)
- overcome debootstrap limitation → install a fully upgraded system (multistrap?)

instance creation with ganeti-instance-debootstrap – who can do faster?

```
Terminal - root@gisu828: ~
root@gisu828:~# date && gnt-instance add -t drbd --disk 0:size=1G --net 0:network=gruen97,ip=172.28.97.234 -H kvm:kernel_path=/vmlinuz,1
nitrd_path=/initrd.img,root_path=/dev/vda1,kernel_args='ro elevator=noop net.ifnames=0' -B vcpus=1,memory=1G -o debootstrap+default --no
-name-check --no-ip-check test.vm
Sun Jun 16 11:06:49 CEST 2019
Sun Jun 16 11:06:53 2019 - INFO: Selected nodes for instance test.vm via iallocator hail: gisu850.gisa-halle.de, gisu851.gisa-halle.de
Sun Jun 16 11:06:53 2019 - INFO: NIC/0 inherits netparams ['br97', 'bridged', u'']
Sun Jun 16 11:06:55 2019 * creating instance disks...
Sun Jun 16 11:06:57 2019 adding instance test.vm to cluster config
Sun Jun 16 11:06:57 2019 adding disks to cluster config
Sun Jun 16 11:06:58 2019 * wiping instance disks...
Sun Jun 16 11:06:58 2019 - INFO: * Wiping disk 0
Sun Jun 16 11:06:58 2019 - INFO: - done: 10.0% ETA: 3s
Sun Jun 16 11:07:01 2019 - INFO: Waiting for instance test.vm to sync disks
Sun Jun 16 11:07:01 2019 - INFO: Instance test.vm's disks are in sync
Sun Jun 16 11:07:01 2019 - INFO: Waiting for instance test.vm to sync disks
Sun Jun 16 11:07:01 2019 - INFO: Instance test.vm's disks are in sync
Sun Jun 16 11:07:01 2019 * running the instance OS create scripts...
Sun Jun 16 11:07:03 2019 * starting instance...
root@gisu828:~# █

Terminal - giscon3: ~
@giscon3:~$ ping 172.28.97.234 | perl -nle 'print scalar(localtime), " ", $_'
Sun Jun 16 11:07:19 2019 PING 172.28.97.234 (172.28.97.234) 56(84) bytes of data.
Sun Jun 16 11:07:19 2019 64 bytes from 172.28.97.234: icmp_seq=126 ttl=60 time=1.85 ms
^C
```

THANKS

Questions?